

Gaussian Mixture Model and EM (Expectation Maximization) Algorithm



Changshui Zhang
Dept. of Automation
Tsinghua University

zcs@mail.tsinghua.edu.cn



Reference

- 新教材电子版
- Jeff A. Bilmes, A Gentle Tutorial of the Algorithm and its Application to Parameter Estimation for Gaussian Mixture and Hidden Markov Models



GMM (Gaussian Mixture Model)

- Observed data are N samples independently generated from the following probabilistic model

$$P(X|\Theta) = \sum_{i=1}^M \alpha_i p_i(X|\theta_i)$$

where $\Theta = (\alpha_1, \dots, \alpha_M, \theta_1, \dots, \theta_M)$ and $\sum_{i=1}^M \alpha_i = 1$

$$p(x|\omega_j, \theta_j) \sim N(\mu_i, \Sigma_i)$$



GMM (Gaussian Mixture Model)

- Likelihood Function and log Likelihood Function

$$p(X|\theta) = \prod_{i=1}^N \sum_{j=1}^c N(x_i|\mu_j, \Sigma_j) P(\omega_j)$$

$$\ln p(X|\theta) = \sum_{i=1}^N \ln \sum_{j=1}^c N(x_i|\mu_j, \Sigma_j) P(\omega_j)$$

$$P(X|\Theta) = \sum_{i=1}^M \alpha_i p_i(X|\theta_i)$$



GMM (Gaussian Mixture Model)

$$\ln p(X|\theta) = \sum_{i=1}^N \ln \sum_{j=1}^c N(x_i|\mu_j, \Sigma_j) P(\omega_j)$$

$$\frac{\partial \ln p(X|\theta)}{\partial \mu_k} = 0$$

$$\sum_{i=1}^N \frac{N(x_i|\mu_k, \Sigma_k) P(\omega_k)}{\sum_{j=1}^c N(x_i|\mu_j, \Sigma_j) P(\omega_j)} \Sigma_k^{-1} (x_i - \mu_k) = 0$$



GMM (Gaussian Mixture Model)

$$\sum_{i=1}^N \frac{N(x_i | \mu_k, \Sigma_k) P(\omega_k)}{\sum_{j=1}^c N(x_i | \mu_j, \Sigma_j) P(\omega_j)} \Sigma_k^{-1} (x_i - \mu_k) = 0$$

$$P(\omega_k | x_i, \mu_k, \Sigma_k) = \frac{N(x_i | \mu_k, \Sigma_k) P(\omega_k)}{\sum_{j=1}^c N(x_i | \mu_j, \Sigma_j) P(\omega_j)}$$

$$\mu_k = \frac{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k) x_i}{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)}$$



GMM (Gaussian Mixture Model)

$$P(\omega_k | x_i, \mu_k, \Sigma_k) = \frac{N(x_i | \mu_k, \Sigma_k) P(\omega_k)}{\sum_{j=1}^c N(x_i | \mu_j, \Sigma_j) P(\omega_j)}$$

$$\mu_k = \frac{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k) x_i}{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)}$$

$$P(\omega_k) = \frac{1}{N} \sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)$$

$$\Sigma_k = \frac{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k) (x_i - \mu_k)(x_i - \mu_k)^T}{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)}$$



Parameter estimation for GMM

- Initialize: $\mu_k^{new}, \Sigma_k^{new}, P^{new}(\omega_k)$
- E step: $\mu_k = \mu_k^{new}, \Sigma_k = \Sigma_k^{new}, P(\omega_k) = P^{new}(\omega_k)$

$$P(\omega_k | x_i, \mu_k, \Sigma_k) = \frac{N(x_i | \mu_k, \Sigma_k) P(\omega_k)}{\sum_{j=1}^c N(x_i | \mu_j, \Sigma_j) P(\omega_j)}$$

- M step: $P^{new}(\omega_k) = \frac{1}{N} \sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)$

$$\mu_k^{new} = \frac{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k) x_i}{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)}$$

$$\Sigma_k^{new} = \frac{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k) (x_i - \mu_k)(x_i - \mu_k)^T}{\sum_{i=1}^N P(\omega_k | x_i, \mu_k, \Sigma_k)}$$



Maximum Likelihood

- Problem:

The set of parameters: Θ

A data set: $X = \{x_1, x_2, \dots, x_N\}$

A density function: $p(x|\Theta)$

Likelihood function:

$$L(\Theta|X) := p(X|\Theta) = \prod_{i=1}^N p(x_i|\Theta)$$

where the data is fixed.

Task: find $\Theta^* = \arg \max_{\Theta} L(\Theta|X)$



Incomplete data

- Missing values
 - Empty items in forms
 - Forgotten data
- Hidden variables
 - Can not be measured and observed.

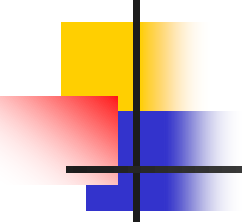


Basic EM

- An elaborate technique of finding the **maximum likelihood estimate** of the parameters of a distribution from a given data set when the data is **incomplete or has missing values**.
- Reference: A.P.Dempster, N.M.Laird, and D.B.Rubin. Maximum-likelihood from incomplete data via the em algorithm. J.Royal Statist. Soc. Ser. B., 39, 1977

- 
- Two main applications of the EM algorithm:
-

- The data indeed has **missing values**, due to problems with or limitations of the observation process
- The optimizing the likelihood function is analytically intractable but when the likelihood function can be **simplified by assuming the existence** of and values for additional but missing (or hidden) parameters.


- 
-
- We assume that data $X = \{x_1, x_2, \dots, x_N\}$ is observed and incomplete data.
 - A complete data set: $Z = (X, Y)$
 - A joint density function:
$$P(Z|\Theta) = p(X, Y|\Theta) = p(Y|X, \Theta)p(X|\Theta)$$

- 
- New likelihood function: complete data likelihood

$$L(\Theta|X, Y) = L(\Theta|X, Y) = p(X, Y|\Theta)$$

- It is a random variable since Θ, X are constant and Y is a random variable.

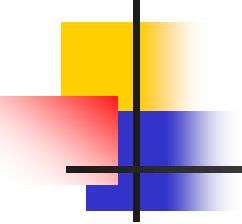
$$L(\theta) = \log p(X|\theta) = \log \left(\sum_Y p(X, Y|\theta) \right)$$



Maximizing the log likelihood directly is often difficult because the log of the sum can potentially couple all of the parameters of the model.

Any distribution $q(y)$ over the hidden variables defines a lower bound on L :

$$L(\theta) \geq \sum_y q(y) \log p(x, y | \theta) - \sum_y q(y) \log q(y) = F(q, \theta)$$



$$\begin{aligned} L(\theta) &= \log \sum_y p(x, y|\theta) = \log \sum_y q(y) \frac{p(x, y|\theta)}{q(y)} \\ &\geq \sum_y q(y) \log \frac{p(x, y|\theta)}{q(y)} \\ &= \sum_y q(y) \log p(x, y|\theta) - \sum_y q(y) \log q(y) \\ &= F(q, \theta) \end{aligned}$$

The inequality is known as **Jensen's inequality**



The EM Algorithm

The Expectation-Maximization algorithm alternates between maximizing F with respect to q and θ , respectively, holding the other fixed.

E step: $q_{[k+1]} \leftarrow \arg \max_q F(q, \theta_{[k]})$

M step: $\theta_{[k+1]} \leftarrow \arg \max_{\theta} F(q_{[k+1]}, \theta)$



The EM Algorithm

The maximum in the **E step** is obtained by setting: $q_{[k+1]}(y) = p(y|x, \theta_{[k]})$

$$F(q, \theta) = \sum_y q(y) \log \frac{p(x, y|\theta)}{q(y)}$$

$$= \sum_y p(y|x, \theta) \log \frac{p(y|x, \theta)p(x|\theta)}{p(y|x, \theta)}$$

$$= \sum_y p(y|x, \theta) \log p(x|\theta)$$

$$= \log p(x|\theta) = L(\theta) \qquad F(q, \theta) \leq L(\theta)$$



The EM Algorithm: E Step

$$q_{[k+1]}(y) = p(y|x, \theta_{[k]})$$

$$F(q_{[k+1]}, \theta) = \sum_y q_{[k+1]}(y) \log p(x, y|\theta) \\ - \sum_y q_{[k+1]}(y) \log q_{[k+1]}(y)$$

$$Q(\theta_{[k]}, \theta) = \sum_y p(y|x, \theta_{[k]}) \log p(x, y|\theta) \\ = E[\ln p(x, y|\theta) | x, \theta_{[k]}]$$

The EM Algorithm: M Step

The maximum in the **M step** is obtained by maximizing the first term of $F(q, \theta)$

M step: $\theta_{[k+1]} \leftarrow \arg \max_{\theta} \sum_y p(y|x, \theta_{[k]}) \log p(y, x|\theta)$

$$\begin{aligned} Q(\theta_{[k]}, \theta) &= \sum_y p(y|x, \theta_{[k]}) \log p(x, y|\theta) \\ &= E[\ln p(x, y|\theta) | x, \theta_{[k]}] \end{aligned}$$



EM Algorithm


- Initialize $\theta_{[0]}$ with random/guess, set $i=1$
- E-step:

$$\begin{aligned} Q(\theta_{[k]}, \theta) &= \sum_y p(y|x, \theta_{[k]}) \log p(x, y|\theta) \\ &= E[\ln p(x, y|\theta) | x, \theta_{[k]}] \end{aligned}$$

- M-step: $\theta_{[k+1]} \leftarrow \arg \max_{\theta} Q(\theta_{[k]}, \theta)$
- $i = i+1$
- repeat until convergence

EM Algorithm()

- The figure is Copied from book by Bishop





Generalized EM

- Assume $\ln p(X|\theta)$ and Q function are differentiable in θ . The EM likelihood converges to a point where $\frac{\partial}{\partial \theta} \ln p(X|\theta) = 0$
- GEM: Instead of setting $\theta_{[k+1]} = \arg \max_{\theta} Q(\theta_{[k]}, \theta)$
Just find $\theta(n)$ such that
$$Q(\theta_{[k]}, \theta_{[k+1]}) > Q(\theta_{[k]}, \theta)$$
- GEM also is guaranteed to converge



Finding Maximum Likelihood Mixture Densities Parameters via EM

- Observed data are N samples independently generated from the following probabilistic model

$$p(X|\Theta) = \sum_{i=1}^M \alpha_i p_i(X|\theta_i)$$

where $\Theta = (\alpha_1, \dots, \alpha_M, \theta_1, \dots, \theta_M)$ and $\sum_{i=1}^M \alpha_i = 1$



Finding Maximum Likelihood Mixture Densities Parameters via EM

- The unknown parameters $y_i \in 1, \dots, M$ indicate that the i^{th} sample is generated by the y_i^{th} mixture component
- Though $\ln \prod_{i=1}^N p(x_i | \Theta) = \sum_{i=1}^N \ln(\sum_{j=1}^M \alpha_j p_j(x_i | \theta_j))$ is hard to compute, we have a much easier choice:

$$\begin{aligned} \ln(p(X, Y | \Theta)) &= \sum_{i=1}^N \ln(P(x_i | y_i) P(y_i)) \\ &= \sum_{i=1}^N \ln(\alpha_{y_i} p_{y_i}(x_i | \theta_{y_i})) \end{aligned}$$



E-step

- Guess an initial $\Theta^g = (\alpha_1^g, \dots, \alpha_M^g, \theta_1^g, \dots, \theta_M^g)$
- Expectation equation

$$\begin{aligned} Q(\Theta, \Theta^g) &= \sum_Y \ln(p(X, Y|\Theta)) p(Y|X, \Theta^g) \\ &= \sum_y \sum_{i=1}^N \ln(\alpha_{y_i} p_{y_i}(x_i|\theta_{y_i})) \prod_{j=1}^N p(y_j|x_j, \Theta^g) \end{aligned}$$

$$Q(\theta_{[k]}, \theta) = \sum_y p(y|x, \theta_{[k]}) \log p(x, y|\theta)$$



E-step

$$\begin{aligned}
 & \sum_y \sum_{i=1}^N \ln(\alpha_{y_i} p_{y_i}(x_i | \theta_{y_i})) \prod_{j=1}^N p(y_j | x_j, \Theta^g) \\
 &= \sum_{y_1=1}^M \sum_{y_2=1}^M \dots \sum_{y_N=1}^M \sum_{i=1}^N \ln(\alpha_{y_i} p_{y_i}(x_i | \theta_{y_i})) \prod_{j=1}^N p(y_j | x_j, \Theta^g) \\
 &= \sum_{y_1=1}^M \sum_{y_2=1}^M \dots \sum_{y_N=1}^M \sum_{i=1}^N \sum_{l=1}^M \delta_{l,y_i} \ln(\alpha_l p_l(x_i | \theta_l)) \prod_{j=1}^N p(y_j | x_j, \Theta^g) \\
 &= \sum_{l=1}^M \sum_{i=1}^N \ln(\alpha_l p_l(x_i | \theta_l)) \sum_{y_1=1}^M \sum_{y_2=1}^M \dots \sum_{y_N=1}^M \delta_{l,y_i} \prod_{j=1}^N p(y_j | x_j, \Theta^g)
 \end{aligned}$$



E-step

$$\sum_{y_1=1}^M \sum_{y_2=1}^M \cdots \sum_{y_N=1}^M \delta_{l,y_i} \prod_{j=1}^N p(y_j|x_j, \Theta^g)$$

$$= \left(\sum_{y_1=1}^M \cdots \sum_{y_{i-1}=1}^M \sum_{y_{i+1}=1}^M \cdots \sum_{y_N=1}^M \prod_{j=1, j \neq i}^N p(y_j|x_j, \Theta^g) \right) p(l|x_i, \Theta^g)$$

$$= \prod_{j=1, j \neq i}^N \left(\sum_{y_j=1}^M p(y_j|x_j, \Theta^g) \right) p(l|x_i, \Theta^g)$$

$$= p(l|x_i, \Theta^g)$$



E-step

$$\begin{aligned} Q(\Theta, \Theta^g) &= \sum_{l=1}^M \sum_{i=1}^N \ln(\alpha_l p_l(x_i | \theta_l)) p(l | x_i, \Theta^g) \\ &= \sum_{l=1}^M \sum_{i=1}^N \ln(\alpha_l) p(l | x_i, \Theta^g) \\ &\quad + \sum_{l=1}^M \sum_{i=1}^N \ln(p_l(x_i | \theta_l)) p(l | x_i, \Theta^g) \end{aligned}$$

Note: The term containing θ_l and the term containing α_l are not related and can be maximized independently



M-step

- Maximize α_l : Introduce Lagrange multiplier λ

$$\frac{\partial}{\partial \alpha_l} [Q(\Theta, \Theta^g) + \lambda(\sum_l \alpha_l - 1)] = 0$$

$$\frac{\partial}{\partial \alpha_l} \left[\sum_{l=1}^M \sum_{i=1}^N \ln(\alpha_l) p(l|x_i, \Theta^g) + \lambda(\sum_l \alpha_l - 1) \right] = 0$$

$$\sum_{i=1}^N \frac{1}{\alpha_l} p(l|x_i, \Theta^g) + \lambda = 0$$

$$\sum_{i=1}^N p(l|x_i, \Theta^g) + \lambda \alpha_l = 0 \quad l = 1, \dots, M$$



M-step

$$\sum_{i=1}^N p(l|x_i, \Theta^g) + \lambda \alpha_l = 0 \quad l = 1, \dots, M$$

$$\sum_{l=1}^M \sum_{i=1}^N p(l|x_i, \Theta^g) + \sum_{l=1}^M \lambda \alpha_l = 0$$

$$N + \lambda = 0$$

$$\alpha_l = \frac{1}{N} \sum_{i=1}^N p(l|x_i, \Theta^g)$$



M-step

- Maximize θ_l : If $p(x|\theta)$ is d -dimensional Gaussian distributions with mean μ and covariance matrix Σ ,
we can have an analytical expressions for μ and Σ .



M-step

$$B = \sum_{l=1}^M \sum_{i=1}^N \ln(p_l(x_i|\theta_l)) p(l|x_i, \Theta^g)$$

$$= \sum_{l=1}^M \sum_{i=1}^N \left(-\frac{\ln(|\Sigma_l|)}{2} - \frac{(x_i - \mu_l)^T \Sigma_l^{-1} (x_i - \mu_l)}{2} \right) p(l|x_i, \Theta^g)$$

$$\frac{\partial B}{\partial \mu_l} = 0, \quad \sum_{i=1}^N \Sigma_l^{-1} (x_i - \mu_l) p(l|x_i, \Theta^g) = 0$$

$$\frac{\partial x^T A x}{\partial x} = (A + A^T)x, \quad \mu_l = \frac{\sum_{i=1}^N x_i p(l|x_i, \Theta^g)}{\sum_{i=1}^N p(l|x_i, \Theta^g)}$$



M-step

$$\begin{aligned} B &= \sum_{l=1}^M \sum_{i=1}^N \left(-\frac{\ln(|\Sigma_l|)}{2} - \frac{(x_i - \mu_l)^T \Sigma_l^{-1} (x_i - \mu_l)}{2} \right) p(l|x_i, \Theta^g) \\ &= \sum_{l=1}^M \left[\frac{\ln(|\Sigma_l^{-1}|)}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) - \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) \operatorname{tr}(\Sigma_l^{-1} (x_i - \mu_l)(x_i - \mu_l)^T) \right] \\ &= \sum_{l=1}^M \left[\frac{\ln(|\Sigma_l^{-1}|)}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) - \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) \operatorname{tr}(\Sigma_l^{-1} N_{l,i}) \right] \\ N_{l,i} &= (x_i - \mu_l)(x_i - \mu_l)^T \end{aligned}$$



M-step

$$\frac{\partial B}{\partial \Sigma_l^{-1}} = 0$$

$$\begin{aligned} B &= \sum_{l=1}^M \left[\frac{\ln(|\Sigma_l^{-1}|)}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) - \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) \text{tr}(\Sigma_l^{-1} N_{l,i}) \right] \\ &\quad - \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) (2\Sigma_l - \text{diag}(\Sigma_l)) - \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) (2N_{l,i} - \text{diag}(N_{l,i})) \\ &= \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) (2M_{l,i} - \text{diag}(M_{l,i})) \\ &= 2S - \text{diag}(S), \quad M_{l,i} = \Sigma_l - N_{l,i}, \quad S = \frac{1}{2} \sum_{i=1}^N p(l|x_i, \Theta^g) M_{l,i} \end{aligned}$$



M-step

$$2S - \text{diag}(S) = 0$$

$$S = 0$$

$$S = \frac{1}{2} \sum_i^N p(l|x_i, \Theta^g) (\Sigma_l - N_{l,i}) = 0$$

$$\Sigma_l = \frac{\sum_{i=1}^N p(l|x_i, \Theta^g) N_{l,i}}{\sum_{i=1}^N p(l|x_i, \Theta^g)} = \frac{\sum_{i=1}^N p(l|x_i, \Theta^g) (x_i - \mu_l)(x_i - \mu_l)^T}{\sum_{i=1}^N p(l|x_i, \Theta^g)}$$

$$\frac{\partial \ln|A|}{\partial A} = 2A^{-1} - \text{diag}(A^{-1})$$

$$\frac{\partial \text{tr}(AB)}{\partial A} = B + B^T - \text{diag}(B)$$



M-step


$$\alpha_l^{new} = \frac{1}{N} \sum_{i=1}^N p(l|x_i, \Theta^g)$$

$$\mu_l^{new} = \frac{\sum_{i=1}^N x_i p(l|x_i, \Theta^g)}{\sum_{i=1}^N p(l|x_i, \Theta^g)}$$

$$\Sigma_l^{new} = \frac{\sum_{i=1}^N p(l|x_i, \Theta^g) (x_i - \mu_l^{new})(x_i - \mu_l^{new})^T}{\sum_{i=1}^N p(l|x_i, \Theta^g)}$$

Experiments


- 5 points from a Gaussian Distribution(3 times)
- Blue: original distribution
- Red: estimated distribution
- Green: samples



(a) $\mu = 0.5910, \sigma^2 = 1.9690$ (b) $\mu = 0.5136, \sigma^2 = 0.5467$ (c) $\mu = 0.2471, \sigma^2 = 0.3051$

Experiments

- 50 points from a Gaussian Distribution (3 times)
- Blue: original distribution
- Red: estimated distribution
- Green: samples




(d) $\mu = 0.0560, \sigma^2 = 1.1030$ (e) $\mu = 0.1898, \sigma^2 = 1.1726$ (f) $\mu = 0.2174, \sigma^2 = 1.3243$

Experiments

- 500,5000,50000 points from a Gaussian Distribution

- Blue: original distribution




(g) $\mu = 0.0247, \sigma^2 = 1.0581$ (h) $\mu = -0.0095, \sigma^2 = 1.0078$ (i) $\mu = 0.0107, \sigma^2 = 1.0064$

Experiments

■ 500 points from a uniform Distribution


- Blue: original distribution
- Red: estimated distribution
- Green: samples



Experiments

■ 5000 points from a GMM


- Blue: original distribution
- Red: estimated distribution
- Green: samples



(3-5) 混合高斯分布样本的正态分布最大似然估计(5000 样本)


Experiments

- 5 points from a GMM (2 times)
- Blue: original distribution
- Red: estimated distribution
- Green: samples



Experiments

- 50 points from a GMM (2 times)
- Blue: original distribution
- Red: estimated distribution
- Green: samples



Experiment

- 500 and 5000 points from a GMM (2 times)
- Blue: original distribution
- Red: estimated distribution
- Green: samples




图 3-8. 一维 500 个样本的 EM 算法处理结果, 蓝线为采样概率密度函数, 红线为 EM 估计概率密度函数






图 3-9. 一维 5000 个样本的 EM 算法处理结果, 蓝线为采样概率密度函数, 红线为 EM 估计概率密度函数


Experiments

- A wrong number of components is given:




Experiments

- A wrong number of components is given:




(a)



(b)

Experiments

- Overfitting



- The figure is Copied from book by Duda



Research

- Parzon window on a manifold
- Stochastic EM
- Convergence of EM
- **<http://www.vision.caltech.edu/welling/class/LearningSystemsB.html>**



Research


- **PCA**
- **ICA**
- **Factor Analysis**
- **K-means**
- **Mixture of Gaussians**
- **Generative Topographic Mapping**



Research

- **Cluster Weighted Models**
- **Mixture of Experts**
- **Kalman Filter**
- **HMM**
- **Helmholtz Machine**
- **Boltzman Machine**
- **It will be shown that most of the learning schemes for these models can be understood as versions of the Expectation Maximization (EM) algorithm, thus providing a unified view.**

applications





Research

- How many components in GMM?
- MDL (Minimum description length) method
- Competitive EM Algorithm for Finite Mixture Models
- Merge and Split operation and RJMCMC
- Kernel GMM



Research

EM:

- Machine Learning
- Computer Vision
- Pattern Recognition
- Bayesian Networks



Research

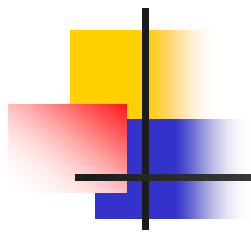
GMM

- background modeling in video tracking
- Speaker identification
- ...



Reference:

- J. A. Bilmes et al "A Gentle Tutorial of the EM Algorithm and its application to Parameter Estimation for Gaussian Mixture and Hidden Markov Models" 1998
- A.P.Dempster et al "Maximum-likelihood from incomplete data via the EM algorithm" 1977
- T.K.Moon "The Expectation-Maximization Algorithm" *IEEE Trans Signal Processing* 1996



Thanks!